Cognitive Physics as a Constraint-Based Meeting Ground

Persistence, Learning, and Stability Under Uncertainty

Joel Peña Muñoz Jr.

January 27, 2026

Abstract

The sciences of cognition are fragmented across disciplines that differ in scale, formalism, and explanatory language, including neuroscience, psychology, artificial intelligence, control theory, evolutionary biology, information theory, and philosophy of mind. Attempts at unification often fail by either reducing one domain to another or proposing a final theory that prematurely closes inquiry. This paper proposes *Cognitive Physics* as an alternative: a constraint-based meeting ground in which diverse models of cognition can coexist, compete, and evolve, provided they satisfy non-negotiable requirements of persistence under uncertainty. Rather than unifying cognition by shared representations or mechanisms, Cognitive Physics unifies it through structural constraints on stability, aggregation, and failure. Cognition is treated not as a privileged substance or top-down causal agent, but as a stable, resilient pattern that emerges when interacting components remain reconstructable across scales. This framework provides a common language for comparing models without erasing their differences, enabling continuous learning, theoretical evolution, and cross-disciplinary integration without invoking a final theory of mind.

Figure 1: Cognitive Physics as a constraint-based meeting ground: diverse models coexist as long as they remain viable under uncertainty.

1 Introduction

Cognition is studied everywhere and agreed upon almost nowhere. Neural firing patterns, symbolic manipulation, predictive processing, dynamical systems, evolutionary adaptation, embodied interaction, and artificial learning architectures all claim partial authority over what it means to think, learn, or decide. Each framework is locally successful, yet none comfortably subsumes the others without distortion.

This fragmentation is not accidental. Cognitive phenomena span multiple spatial, temporal, and organizational scales, from millisecond neural events to lifelong learning, from molecular signaling to cultural evolution. As a result, attempts to impose a single explanatory language often collapse under their own ambition, either by oversimplifying cognition or by drifting into unfalsifiable abstraction.

Cognitive Physics begins from a different premise: cognition does not require a single model, but it does require survival. Any system that learns, adapts, or behaves intelligently must remain viable under uncertainty. This simple observation introduces a powerful organizing principle. Instead of asking which model of cognition is *true*, Cognitive Physics asks which models can *persist* when subjected to disturbance, noise, novelty, and change.

The result is not a final theory, but a shared constraint space—a meeting ground where multiple sciences of cognition can enter, interact, and be compared without being reduced to one another.

2 Why Unification Has Failed

Historically, unification efforts in cognitive science have followed one of two paths. The first is reductionism: cognition is explained entirely in terms of neural mechanisms, computational primitives, or physical dynamics. The second is abstraction: cognition is treated as a high-level phenomenon independent of physical or biological constraints.

Both approaches encounter predictable limits. Reductionism struggles with scale and emergence, while abstraction struggles with grounding and falsifiability. Cognitive Physics rejects both extremes by separating *structural constraints* from *representational commitments*. Models are not required to agree on what cognition is made of, only on what cognition must withstand.

Figure 2: Two common failure modes of unification: reductionism (collapsing levels into one) and abstraction (floating free of physical constraints). Cognitive Physics unifies by constraints instead.

3 Cognitive Physics: The Party Analogy (Informal Motivation)

Cognitive Physics can be understood informally as a large, open gathering of ideas—a shared dance floor where very different models of cognition are welcome to show up, mingle, compete, and evolve. There is no dress code, no preferred style, and no central authority deciding which approach is correct. What matters is whether a model can keep its balance when the environment becomes unpredictable.

Some models dance with symbols, others with neurons, equations, feedback loops, or evolutionary histories. They do not need to move the same way. They only need to stay upright when the rhythm changes. Those that adapt, recover, and continue belong on the floor. Those that fall apart are not rejected by taste or fashion; they simply cannot keep dancing under the constraints imposed by reality.

This metaphor captures the spirit of the framework, but the remainder of this paper replaces metaphor with formal structure.

4 Core Principle: Persistence Under Uncertainty

At the heart of Cognitive Physics is a single organizing principle:

> *Any cognitive system must maintain coherence while absorbing uncertainty, or it ceases to function as a cognitive system.*

This principle does not specify mechanisms, architectures, or representations. Instead, it defines admissibility. A model of cognition is admissible if it can explain how learning, adaptation, or behavior remains stable under perturbation, and how failure occurs when those conditions are violated.

From this principle follow four structural conditions that govern entry into the Cognitive Physics framework.

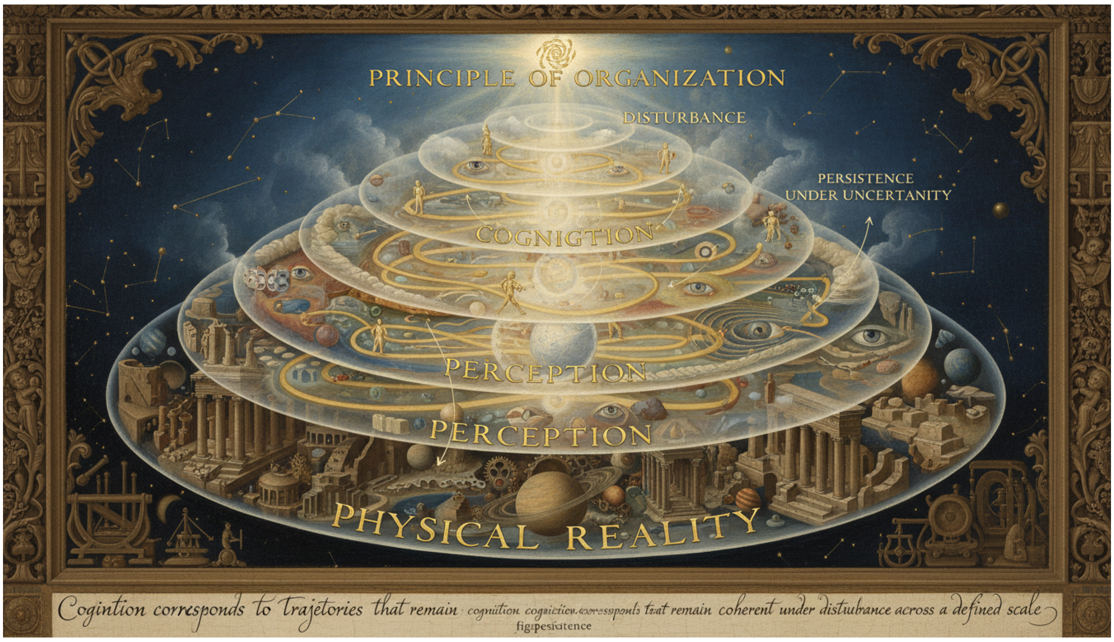

Figure 3: Persistence under uncertainty as the organizing principle: cognition corresponds to trajectories that remain coherent under disturbance across a defined scale.

5 The Four Structural Conditions

5.1 Bounded Disturbance

A cognitive system must prevent perturbations from amplifying without bound. Noise, novelty, or error may disrupt the system, but they cannot cascade indefinitely without destroying coherence. Models that lack mechanisms for damping, regulation, or constraint fail this condition.

5.2 Scale-Indexed Coherence

Cognitive descriptions must be indexed to a specific observation scale. Claims that assume scale-invariant cognition obscure the aggregation processes that produce stability. A valid model must specify the level at which its variables, dynamics, and coherence claims apply.

5.3 Connectivity Viability

The internal interactions of a cognitive system must support persistent macroscopic patterns. Fragmented or weakly connected components may compute locally, but they cannot sustain cognition globally. Viable cognition emerges from dense, structured connectivity that preserves reconstructability over time.

5.4 Collapse Accountability

Every admissible model must specify its own failure modes. A theory that cannot explain how cognition degrades, collapses, or loses coherence under violated constraints is incomplete. Collapse is not an embarrassment; it is a diagnostic.

6 Cognition as a Stability Basin

Within Cognitive Physics, cognition is not treated as a top-down controller, an inner narrator, or a privileged causal agent. Instead, it is understood as a stability basin: a region of state space in which system trajectories remain bounded under admissible perturbations.

Learning corresponds to movement within or between such basins. Adaptation corresponds to basin reshaping. Identity corresponds to basin persistence. None of these require invoking metaphysical agency or unexplained initiation.

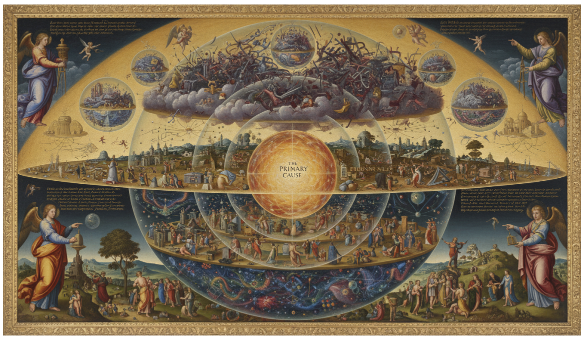

(a) Bounded disturbance.

(b) Scale-indexed coherence.

(c) Connectivity viability.

(d) Collapse accountability.

Figure 4: The four structural entry conditions for models inside the Cognitive Physics meeting ground.

7 Roadmap: Integrating the Sciences of Cognition

The remainder of this paper will systematically integrate major cognitive frameworks into the Cognitive Physics constraint space, including:

- Neuroscience and neural dynamics
- Psychology and behavioral regularities
- Bayesian inference and predictive processing
- Control theory and stability analysis
- Information theory and entropy regulation
- Evolutionary biology and adaptive landscapes
- Embodied and enactive cognition
- Artificial intelligence and machine learning
- Ethics as viability under constraint

- Social and collective cognition

Each framework will be evaluated not for correctness, but for viability: how it satisfies the four structural conditions, where it excels, and where it fails.

8 Neuroscience Within the Cognitive Physics Framework

Neuroscience is often treated as the foundational science of cognition, on the assumption that understanding neural mechanisms will eventually explain all cognitive phenomena. While neural dynamics are indispensable, Cognitive Physics reframes their role. Neuroscience does not supply the definition of cognition; it supplies one class of systems that successfully satisfies the constraints required for cognition to persist.

From the perspective of Cognitive Physics, the central question is not whether neurons cause cognition, but how neural systems maintain coherence under constant uncertainty, noise, and perturbation. The human brain is not a precision instrument operating in isolation. It is an energetically constrained, thermally noisy, metabolically fragile system embedded in an unpredictable environment. That it supports stable perception, learning, and behavior at all is therefore a remarkable fact requiring structural explanation.

8.1 Neural Activity as Event Aggregation

At microscopic scales, neural systems are composed of discrete events: ion channel openings, synaptic releases, action potentials, and molecular signaling cascades. These events occur stochastically, are sensitive to noise, and are individually insufficient to explain macroscopic cognitive stability.

Cognitive Physics treats these neural events as elements of a larger aggregation process. Cognition does not reside in any single spike or synapse, but in the stable patterns that emerge when vast numbers of such events are integrated over appropriate spatial and temporal windows. This perspective aligns naturally with empirical findings in population coding, distributed representation, and ensemble dynamics.

The requirement of scale-indexed coherence is particularly salient here. Neural explanations that oscillate ambiguously between single-neuron causation and system-level behavior often fail to specify the observation scale at which claims are made. Cognitive Physics insists on this specification, treating neural descriptions as effective models whose validity is indexed to defined aggregation windows.

8.2 Bounded Disturbance in Neural Systems

Neural systems are continuously perturbed by internal and external sources of variability: thermal noise, synaptic fluctuations, sensory ambiguity, and metabolic constraints. Despite this, cognition remains remarkably stable. This stability is not accidental; it is enforced through multiple overlapping mechanisms that satisfy the bounded disturbance condition.

Recurrent connectivity, inhibitory-excitatory balance, homeostatic plasticity, and neuromodulatory regulation all function to prevent perturbations from amplifying uncontrollably. When these mechanisms fail, cognition degrades in predictable ways, manifesting as seizures, hallucinations, attentional collapse, or loss of consciousness.

From the Cognitive Physics perspective, such pathologies are not anomalies but confirmations. They reveal the boundaries of the stability basin within which cognition persists. Neuroscience thus contributes not only mechanisms of function, but clear empirical illustrations of collapse accountability.

8.3 Connectivity Viability and Recurrent Structure

A defining feature of neural cognition is dense, recurrent connectivity. Feedforward architectures alone are insufficient to sustain cognitive stability over time. Instead, feedback loops, lateral interactions, and multi-scale recurrence create conditions under which macroscopic patterns can persist despite continuous micro-level turnover.

Cognitive Physics interprets this recurrence as a structural necessity rather than a design choice. Without viable connectivity, neural activity fragments into transient, uncorrelated events incapable of supporting memory, perception, or action. With sufficient connectivity, the system forms stable basins of activity that can be re-entered, reshaped, or transitioned between under learning.

Importantly, this framing avoids attributing causal privilege to any particular neural structure or region. What matters is not where cognition happens, but whether the overall connectivity supports bounded aggregation.

8.4 Learning as Basin Reshaping

In Cognitive Physics, learning is not the accumulation of information in a static substrate, nor the optimization of a single objective function. It is the gradual reshaping of stability basins under sustained interaction with uncertainty.

Neural plasticity mechanisms—synaptic modification, structural remodeling, neuromodulatory gating—are interpreted as physical processes that alter the geometry of these basins. Learning succeeds when such reshaping enhances persistence under anticipated perturbations. Learning fails when plasticity destabilizes coherence faster than new structure can form.

This view naturally accommodates both rapid adaptation and long-term consolidation without invoking separate cognitive faculties. It also predicts that excessive plasticity can be as destructive as insufficient plasticity, a prediction borne out in both developmental disorders and neurodegenerative disease.

8.5 No Privileged Neural Homunculus

Traditional cognitive explanations often smuggle agency back into neuroscience by attributing control to executive centers, decision modules, or supervisory processes. Cognitive Physics rejects this move. Neural systems do not require a homunculus to coordinate activity; coordination emerges when connectivity and aggregation satisfy stability constraints.

Executive function, attention, and intention are therefore interpreted as macroscopic stability phenomena, not initiating causes. They describe how neural activity remains organized over time, not how it is commanded from above.

This interpretation preserves the empirical successes of neuroscience while eliminating unnecessary metaphysical commitments.

8.6 What Neuroscience Contributes to the Meeting Ground

Within the Cognitive Physics framework, neuroscience contributes:

- Concrete examples of bounded disturbance in action
- Empirical maps of connectivity viability
- Measurable collapse modes under violated constraints

- Physical mechanisms for basin formation and reshaping

At the same time, neuroscience does not close the question of cognition. Its models remain scale-bound, substrate-specific, and incomplete without integration with broader constraint principles.

Cognitive Physics therefore does not replace neuroscience. It hosts it—placing neural models on the shared dance floor alongside other sciences of cognition, where they can be compared, challenged, and extended without being crowned as final arbiters.

9 Bayesian Inference and Predictive Processing

Bayesian inference and predictive processing have become central frameworks in contemporary cognitive science. They propose that cognition consists fundamentally of probabilistic inference: systems generate predictions about incoming signals and update internal models by minimizing prediction error. In many accounts, this process is elevated from a useful formalism to a near-universal explanation of perception, action, and learning.

Cognitive Physics neither rejects nor canonizes Bayesian approaches. Instead, it asks a more basic question: under what conditions does Bayesian inference remain viable as a cognitive process? Framed this way, Bayesian cognition becomes one participant at the meeting ground rather than a final arbiter of cognition itself.

9.1 Bayesian Models as Stability-Preserving Strategies

At its core, Bayesian inference is a method for maintaining coherence under uncertainty. Priors encode accumulated structure, likelihoods encode sensitivity to new information, and posteriors represent updated beliefs that balance both. From the perspective of Cognitive Physics, this is not a metaphysical claim about the nature of mind, but a concrete strategy for satisfying bounded disturbance.

A purely reactive system that updates without priors is unstable. A system that clings rigidly to priors without updating is brittle. Bayesian updating occupies a narrow regime between these extremes, where perturbations are absorbed without erasing structure. In this sense, Bayesian inference can be interpreted as one way—among others—to maintain persistence under uncertainty.

This reframing strips Bayesian cognition of its privileged status while preserving its functional significance.

9.2 Prediction Error as a Disturbance Signal

Predictive processing frameworks emphasize prediction error as the primary driver of learning and perception. Cognitive Physics interprets prediction error not as an objective quantity to be minimized at all costs, but as a structured disturbance signal that must be regulated.

Prediction error is informative only insofar as it remains bounded. Excessive error destabilizes internal models; insufficient error prevents adaptation. Systems that blindly minimize prediction error risk pathological behavior, including hallucination-like overfitting or withdrawal from informative input.

Thus, the relevant question is not whether cognition minimizes error, but whether it maintains error within a viable range. This places predictive processing squarely within the bounded disturbance condition.

9.3 Scale Dependence in Bayesian Cognition

Bayesian models are often presented as scale-invariant, applying equally to neurons, brains, organisms, or societies. Cognitive Physics challenges this assumption. Inference processes are always indexed to an observation scale, whether explicitly acknowledged or not.

At fine scales, stochastic variability dominates and probabilistic descriptions are indispensable. At coarse scales, aggregation suppresses high-frequency uncertainty, and deterministic approximations often suffice. Treating Bayesian inference as universally fundamental obscures the fact that it is an effective description that emerges under particular aggregation conditions.

Cognitive Physics therefore treats Bayesian cognition as scale-indexed: valid within specific resolution windows and liable to breakdown outside them.

9.4 Learning as Model Restructuring, Not Just Updating

Standard Bayesian updating assumes a fixed model structure and updates only parameters or beliefs. However, real cognitive systems frequently undergo deeper transformations: new variables are introduced, old ones discarded, and representational frameworks reorganized.

From the Cognitive Physics perspective, this corresponds to basin reshaping rather than simple posterior updating. Learning is successful when restructuring preserves or enhances stability under future uncertainty. It fails when restructuring destabilizes coherence faster than new structure can compensate.

This distinction explains why purely Bayesian models struggle to account for developmental learning, paradigm shifts, or radical adaptation without auxiliary mechanisms.

9.5 Connectivity Requirements for Predictive Processing

Predictive processing architectures rely heavily on hierarchical and recurrent connectivity. Predictions flow downward, error signals flow upward, and lateral interactions stabilize representations. Without sufficient connectivity, prediction errors fragment and fail to produce coherent updates.

Cognitive Physics interprets this architecture as a specific realization of connectivity viability. The success of predictive processing depends not on its inferential elegance, but on whether its connectivity supports persistent macroscopic patterns under disturbance.

This explains why predictive processing models that ignore physical or architectural constraints often perform well in theory but poorly in implementation.

9.6 Failure Modes and Collapse in Bayesian Systems

Bayesian and predictive processing models exhibit clear collapse modes when constraints are violated. Overconfident priors can lead to delusional stability. Excessive sensitivity to error can lead to instability or noise chasing. Insufficient connectivity can lead to fragmentation and incoherence.

These failures are not exceptions; they are structurally necessary outcomes when bounded disturbance, scale coherence, or connectivity viability are lost. Cognitive Physics treats such failures as essential diagnostic tools rather than anomalies to be explained away.

A Bayesian model that cannot specify its collapse conditions is incomplete by the standards of Cognitive Physics.

9.7 What Bayesian Cognition Contributes to the Meeting Ground

Within the Cognitive Physics framework, Bayesian inference and predictive processing contribute:

- Formal tools for managing uncertainty
- Clear mechanisms for error-regulated adaptation
- Quantitative models of learning under noise
- Explicit trade-offs between stability and flexibility

They do not, however, exhaust cognition. Their viability depends on structural conditions they do not themselves justify. Cognitive Physics provides that justification by situating Bayesian models within a broader landscape of persistence constraints.

Bayesian cognition thus remains a powerful dancer at the party—but not the only one, and not the DJ.

10 Control Theory, Feedback, and Stability

Control theory provides one of the clearest and most physically grounded approaches to understanding how systems remain stable under disturbance. Long before cognition was formalized as computation or inference, control theory addressed a more primitive question: how can a system regulate itself in the presence of noise, delay, and uncertainty?

Cognitive Physics recognizes control theory not as a metaphor for cognition, but as one of its closest structural relatives. Where Bayesian models emphasize inference, control theory emphasizes regulation. Where predictive processing focuses on error signals, control theory focuses on bounded response. Both are different expressions of the same underlying requirement: persistence under uncertainty.

10.1 Stability as the Primary Criterion

In control theory, a system is not judged by whether it reaches an optimal state, but by whether it remains stable when perturbed. Stability is defined relative to a reference condition and is evaluated through bounded trajectories rather than point predictions.

This emphasis aligns directly with Cognitive Physics. Cognition is not defined by correctness, intelligence, or rationality, but by whether coherent behavior persists under disturbance. A perfectly optimal controller that destabilizes under slight perturbation is useless. A suboptimal controller that remains stable is viable.

This reframing immediately clarifies why many cognitively impressive systems fail in real-world environments: they optimize without stabilizing.

10.2 Feedback as a Physical Necessity

Feedback is not an optional design feature in cognitive systems; it is a physical necessity. Open-loop systems may function briefly in predictable environments, but they collapse under sustained uncertainty. Feedback allows a system to sense deviation and counteract it before instability grows.

Cognitive Physics treats feedback as the mechanism by which bounded disturbance is enforced. Whether implemented through neural recurrence, error correction, sensorimotor loops, or adaptive policies, feedback closes the loop between action and consequence.

Importantly, feedback does not require representation or inference. It requires only that system outputs influence future system states in a way that dampens divergence. This makes control theory applicable across biological, mechanical, and artificial substrates.

10.3 Lyapunov Stability and Cognitive Persistence

Lyapunov theory formalizes stability without requiring explicit solutions to system dynamics. Instead, it asks whether trajectories remain bounded near a reference state when perturbed.

From the perspective of Cognitive Physics, Lyapunov stability offers a precise mathematical analogue for cognitive persistence. A cognitive state corresponds to a region of state space within which activity remains bounded. Learning, adaptation, and behavior correspond to motion within or between such regions.

This interpretation avoids treating cognition as a sequence of decisions or symbolic operations. Instead, cognition is understood as the maintenance of stable trajectories in a high-dimensional dynamical system.

10.4 Control Without Central Command

A common misunderstanding of control theory is the assumption that control requires a central controller issuing commands. In reality, many of the most robust control systems are distributed, decentralized, and emergent.

Biological organisms exhibit layered control: reflexes, local feedback loops, and slower adaptive processes coexist without a single point of authority. Cognitive Physics emphasizes this distributed nature, rejecting the idea of a privileged executive homunculus.

Control emerges from structure, not command. Stability arises because interactions are arranged such that deviations are damped, not because a controller chooses to intervene.

10.5 Trade-Offs Between Stability and Flexibility

Control theory makes explicit a trade-off that Cognitive Physics treats as fundamental: increasing stability often reduces flexibility, while increasing responsiveness increases instability.

Highly damped systems resist perturbation but adapt slowly. Highly responsive systems adapt quickly but risk oscillation or collapse. Cognitive systems must operate within a narrow corridor between rigidity and chaos.

This trade-off explains why cognition cannot be perfectly optimal, perfectly rational, or perfectly adaptive. Such ideals ignore the physical cost of stability. Cognitive Physics treats this trade-off not as a limitation to be overcome, but as a defining feature of viable cognition.

10.6 Failure Modes in Control Systems

Control theory is unusually honest about failure. Instability, oscillation, overshoot, and divergence are not surprises; they are expected outcomes when constraints are violated.

These failure modes map cleanly onto cognitive collapse: panic, fixation, hallucination, indecision, or loss of behavioral coherence. From the Cognitive Physics perspective, such phenomena are not mysterious breakdowns of intelligence, but predictable consequences of control limits.

A cognitive model that cannot specify its control failure modes fails the collapse accountability condition.

10.7 What Control Theory Contributes to the Meeting Ground

Within the Cognitive Physics framework, control theory contributes:

- A formal language for bounded disturbance
- Mathematical tools for stability without optimization
- Explicit trade-offs between robustness and adaptability
- Clear and unavoidable failure modes

Control theory does not explain cognition in its entirety. It does not account for representation, meaning, or learning by itself. What it provides is the physical backbone that makes those higher-level phenomena possible.

Cognitive Physics hosts control theory not as a master framework, but as one of the strongest dancers on the floor—setting the rhythm of stability that all others must respect.

11 Information Theory, Entropy, and Reconstructability

Information theory provides the quantitative backbone for understanding uncertainty, compression, and loss in physical and cognitive systems. Originally developed to study communication under noise, it has since become indispensable across neuroscience, artificial intelligence, thermodynamics, and learning theory. Within Cognitive Physics, information theory plays a unifying role by making persistence under uncertainty mathematically explicit.

Rather than treating information as an abstract commodity, Cognitive Physics treats information as a structural property of systems that remain reconstructable under aggregation. Cognition, in this view, is not the accumulation of information per se, but the preservation of usable structure in the presence of entropy.

11.1 Entropy as Loss of Reconstructability

In classical information theory, entropy measures uncertainty or expected surprise. Cognitive Physics adopts a complementary interpretation: entropy quantifies the degree to which fine-grained structure becomes unrecoverable at a given observation scale.

When interactions among components exceed the system's capacity to maintain coherent aggregation, information is not destroyed at the microscopic level but becomes inaccessible to the observer. From this perspective, entropy increase corresponds to diminishing reconstructability rather than disorder in an absolute sense.

This framing aligns naturally with both thermodynamic entropy and cognitive degradation. Forgetting, confusion, and decoherence are all instances of the same structural phenomenon: information has fallen outside the bounds of stable aggregation.

11.2 Bounded Information Flow

Cognitive systems must regulate the rate at which information enters, propagates, and is transformed internally. Unbounded information flow overwhelms processing capacity and destabilizes coherence. Insufficient information flow prevents adaptation.

Information theory formalizes this balance through channel capacity, rate–distortion trade-offs, and coding constraints. Cognitive Physics interprets these not as engineering details, but as necessary conditions for persistence. Learning succeeds when information flow remains within viable bounds; it fails when entropy accumulates faster than structure can be preserved.

This explains why both sensory overload and sensory deprivation impair cognition, despite lying at opposite ends of the information spectrum.

11.3 Compression, Meaning, and Stability

Compression is often treated as an efficiency measure: reducing redundancy while preserving signal. In Cognitive Physics, compression has deeper significance. Compression is how systems retain meaning under uncertainty.

A cognitive system that stores every detail without compression collapses under informational burden. A system that compresses too aggressively loses critical distinctions. Viable cognition occupies a narrow region where compression preserves the distinctions that matter for persistence.

Meaning, in this framework, is not symbolic reference but stability-relevant compression. Patterns are meaningful insofar as they support future coherence under anticipated disturbances.

11.4 Entropy Production and Learning

Learning necessarily produces entropy. Model updates, plasticity, and exploration all introduce temporary instability as old structures are disrupted. Cognitive Physics emphasizes that learning is not entropy minimization, but entropy management.

A system that refuses entropy increase cannot learn. A system that produces entropy without compensatory structure collapses. Successful learning is therefore characterized by transient entropy spikes followed by re-stabilization into new basins of coherence.

This perspective explains why learning is often accompanied by confusion, error, and apparent regression before improvement emerges.

11.5 Information Bottlenecks and Cognitive Viability

Information bottleneck principles formalize the trade-off between relevance and complexity. Cognitive Physics interprets bottlenecks not merely as optimization strategies, but as structural necessities.

Cognitive systems must discard vast amounts of potential information to remain viable. Bottlenecks enforce selectivity, ensuring that retained information supports stability and action rather than overwhelming capacity.

Failures of bottlenecking manifest as rumination, overfitting, paralysis by analysis, or cognitive overload. These are not psychological quirks, but predictable outcomes when entropy regulation fails.

11.6 Scale Dependence of Entropy

Entropy is inherently scale-dependent. Fine-grained descriptions preserve more information but are fragile under noise. Coarse-grained descriptions suppress detail but enhance stability.

Cognitive Physics insists that claims about information processing must specify the scale at which entropy is evaluated. A model that appears information-rich at one scale may be incoherent at another. Conversely, a model that appears lossy may be optimally stable at the scale relevant to behavior.

This resolves long-standing debates between detailed mechanistic explanations and high-level cognitive descriptions by recognizing both as scale-indexed summaries of the same underlying structure.

11.7 Failure Modes: Entropic Collapse

When entropy production exceeds a system's capacity to regulate it, cognitive collapse occurs. This may appear as loss of memory, breakdown of coordination, disintegration of meaning, or failure to act.

Crucially, collapse does not imply destruction of underlying components. Neural activity, symbols, or computations may continue, but without stable aggregation they no longer constitute cognition.

Cognitive Physics treats such collapse as an essential diagnostic. Any theory of cognition that cannot specify its entropic failure modes fails the collapse accountability condition.

11.8 What Information Theory Contributes to the Meeting Ground

Within Cognitive Physics, information theory contributes:

- Quantitative measures of uncertainty and loss
- Formal limits on information flow and storage
- Trade-offs between compression and fidelity
- Clear markers of cognitive overload and collapse

Information theory does not define cognition, but it constrains it. It specifies what no cognitive system can escape: entropy must be managed, not eliminated.

Cognitive Physics hosts information theory as the language that makes uncertainty visible, measurable, and unavoidable—ensuring that the party does not drift into metaphor without math.

12 Evolutionary Cognition and Adaptive Landscapes

Evolutionary theory provides the longest-running and most unforgiving test of cognitive viability. Across millions of years, systems that failed to maintain coherence under environmental uncertainty did not merely malfunction; they disappeared. For this reason, evolutionary cognition occupies a central position within Cognitive Physics. It demonstrates that persistence under uncertainty is not a metaphorical criterion, but a literal selection pressure operating across deep time.

Cognitive Physics does not treat evolution as an external explanation layered on top of cognition. Instead, evolution is understood as the slowest and most comprehensive learning process available to physical systems, shaping which cognitive architectures are even possible.

12.1 Evolution as Constraint Discovery

Evolution does not optimize cognition toward truth, accuracy, or intelligence in any abstract sense. It discovers constraints. Traits persist not because they are correct, but because they support continued existence under fluctuating conditions.

From the perspective of Cognitive Physics, evolution explores the space of possible structures and eliminates those that cannot regulate uncertainty. Cognitive capacities emerge as side effects of increasingly refined stability under disturbance, rather than as explicit targets of selection.

This reframing dissolves the false dichotomy between adaptation and accident. What survives is what fits the constraint landscape, regardless of intention or foresight.

12.2 Adaptive Landscapes as Stability Topographies

Adaptive landscapes are often visualized as fitness surfaces with peaks and valleys. Cognitive Physics reinterprets these landscapes as stability topographies: regions of parameter space where systems remain viable under perturbation.

Fitness peaks correspond to basins of persistence. Valleys correspond to regions where small changes lead to collapse. Movement across the landscape reflects the balance between exploration, exploitation, and survivability.

Importantly, these landscapes are not static. Environmental change reshapes them continuously, ensuring that no configuration remains optimal indefinitely. This dynamic reshaping reinforces the principle that cognition must remain adaptable rather than finalized.

12.3 Evolutionary Trade-Offs and Cognitive Limits

Evolution makes trade-offs unavoidable. Enhancing one cognitive capacity often destabilizes another. Increased sensitivity improves detection but amplifies noise. Enhanced memory preserves structure but reduces flexibility. Faster learning accelerates adaptation but increases error.

Cognitive Physics treats these trade-offs as structural necessities rather than design flaws. Any cognitive model that ignores evolutionary trade-offs risks proposing capacities that cannot persist under real constraints.

These limits explain why cognition appears messy, biased, and imperfect. Perfection is unstable. Viability is not.

12.4 Exploration, Mutation, and Controlled Instability

Mutation introduces instability into otherwise stable systems. From a short-term perspective, this instability appears as error. From a long-term perspective, it is the engine of discovery.

Cognitive Physics interprets exploration—whether genetic, behavioral, or cognitive—as controlled instability. Too little instability traps systems in rigid basins. Too much instability destroys coherence. Evolution tunes this balance over generations.

This principle scales naturally to learning systems, where exploration must be sufficient to discover new basins without overwhelming existing structure.

12.5 Embodied Survival and Environmental Coupling

Evolutionary cognition emphasizes embodiment. Cognitive systems do not evolve in isolation; they are coupled to environments that supply energy, impose constraints, and generate uncertainty.

Cognitive Physics highlights this coupling as essential. Cognition is not a property of brains alone, but of brain–body–environment systems that maintain coherence together. Adaptive success depends as much on environmental fit as on internal processing.

This perspective precludes disembodied cognition models that ignore energetic, material, or ecological constraints.

12.6 Failure, Extinction, and Collapse

Evolution is brutally explicit about collapse. Systems that fail to regulate uncertainty are removed from the population. There is no appeal to explanation after extinction.

This makes evolutionary cognition a powerful validation tool for Cognitive Physics. Any proposed cognitive architecture that could not plausibly survive evolutionary pressures violates the collapse accountability condition.

Extinction, in this framework, is not tragedy but information. It marks the boundaries of viability.

12.7 What Evolutionary Cognition Contributes to the Meeting Ground

Within Cognitive Physics, evolutionary theory contributes:

- A deep-time test of persistence under uncertainty
- Concrete demonstrations of stability–flexibility trade-offs
- Natural mechanisms for exploration and basin discovery
- Unavoidable constraints imposed by embodiment and environment

Evolution does not explain cognition in detail, but it explains why cognition must obey constraints long before it ever becomes conscious, symbolic, or intelligent.

Cognitive Physics hosts evolutionary cognition as the longest-running dancer at the party—one that never stops reminding the others that survival is the ultimate arbiter.

13 Embodied and Enactive Cognition

Embodied and enactive approaches to cognition challenge the assumption that cognition is primarily an internal process occurring inside a brain or computational core. Instead, they emphasize that cognition arises through continuous interaction among brain, body, and environment. From the perspective of Cognitive Physics, this shift is not merely philosophical; it reflects a deeper recognition that stability under uncertainty cannot be achieved by internal computation alone.

Cognitive Physics treats embodiment not as an optional embellishment to cognition, but as a structural contributor to persistence. Bodies, sensors, effectors, and environmental couplings participate directly in regulating disturbance, shaping learning, and sustaining coherent behavior over time.

13.1 Action as Constraint, Not Output

In many classical cognitive models, action is treated as the endpoint of internal processing: cognition computes, then acts. Embodied cognition inverts this relationship. Action is not merely an output; it is a means by which systems actively shape the uncertainty they encounter.

Cognitive Physics formalizes this intuition by treating action as a constraint on future disturbance. By moving, orienting, manipulating, or withdrawing, a system alters the distribution of inputs it must absorb. This reduces informational burden and stabilizes internal dynamics.

From this perspective, perception and action are inseparable. Perception guides action, and action sculpts perception, forming a closed loop that enforces bounded disturbance.

13.2 Sensorimotor Loops and Stability

Embodied cognition emphasizes sensorimotor loops rather than internal representations alone. Cognitive Physics interprets these loops as physical mechanisms for maintaining coherence across time.

Sensorimotor coupling allows systems to detect deviation early and correct it before instability propagates. This coupling distributes cognitive labor across body and environment, reducing the need for centralized control or detailed internal models.

When sensorimotor loops are disrupted—through injury, deprivation, or environmental mismatch—cognition degrades even if internal processing remains intact. This provides strong empirical support for the connectivity viability condition at the level of organism–environment interaction.

13.3 Environmental Scaffolding and Offloaded Stability

Embodied and enactive theories highlight the role of environmental scaffolding: tools, symbols, social structures, and physical affordances that support cognition. Cognitive Physics treats scaffolding as offloaded stability.

Rather than storing all structure internally, cognitive systems exploit reliable regularities in the environment. Written language, spatial layouts, cultural norms, and technological artifacts reduce internal entropy by externalizing memory and constraint.

This perspective dissolves rigid boundaries between internal and external cognition. What matters is not where structure resides, but whether it contributes to persistence under uncertainty.

13.4 Enactive Sense-Making

Enactive cognition emphasizes sense-making: the idea that organisms bring forth meaning through active engagement rather than passive representation. Cognitive Physics reframes sense-making as the emergence of stability-relevant distinctions.

A distinction matters when it contributes to survival, coordination, or coherence. Meaning is therefore not intrinsic to stimuli, but relational and constrained by viability. Systems enact worlds that are meaningful precisely because those worlds support persistence.

This reframing grounds enactivism without drifting into relativism. Not all enacted meanings are viable; those that destabilize the system are eliminated by collapse.

13.5 Scale and the Body

Embodiment introduces additional scales into cognition: muscle dynamics, posture, locomotion, and energetic expenditure. Cognitive Physics insists that cognitive explanations account for these scales explicitly.

High-level planning that ignores bodily constraints often fails in practice. Conversely, low-level reflexes can exhibit remarkable intelligence when embedded in appropriate environmental loops.

Recognizing scale-indexed coherence prevents misattributing cognitive failure to internal processing when the true breakdown occurs at the level of bodily or environmental coupling.

13.6 Failure Modes in Embodied Systems

Embodied systems exhibit distinctive collapse modes. Loss of coordination, fatigue, sensory mismatch, or environmental disruption can destabilize cognition even when neural processing remains functional.

These failures highlight an important lesson of Cognitive Physics: cognition is not located in any single component. Collapse can occur anywhere the constraint network fails—brain, body, or environment.

A cognitive model that ignores embodied failure modes fails the collapse accountability condition.

13.7 What Embodied and Enactive Cognition Contribute to the Meeting Ground

Within the Cognitive Physics framework, embodied and enactive approaches contribute:

- Action as a regulator of uncertainty
- Sensorimotor loops as stability mechanisms
- Environmental scaffolding as offloaded coherence
- Meaning as viability-relevant distinction

They remind the meeting ground that cognition is not confined to abstract inference or internal control. It is a physically situated process that survives by reshaping the world it encounters.

Cognitive Physics hosts embodied cognition as the reminder that the dance floor itself matters: the surface, the space, and the rhythm are part of what keeps everyone moving without falling.

14 Artificial Intelligence and Machine Learning

Artificial intelligence and machine learning provide the most explicit and testable arena for Cognitive Physics. Unlike biological systems, artificial systems are designed, trained, deployed, and observed with unprecedented precision. Their successes and failures make visible the constraints that Cognitive Physics treats as fundamental.

Rather than asking whether machines can think, Cognitive Physics asks a more diagnostic question: under what conditions do artificial systems remain coherent, adaptive, and functional when exposed to sustained uncertainty? The answer reveals which cognitive principles are structurally necessary rather than biologically contingent.

14.1 Learning Systems as Physical Systems

Despite being implemented in software, machine learning systems are physical systems. They consume energy, operate under finite precision, experience noise, and interact with unpredictable environments. Training instability, overfitting, catastrophic forgetting, and distributional shift are not abstract failures; they are manifestations of violated constraints.

Cognitive Physics treats learning algorithms as dynamical systems whose trajectories must remain bounded. A model that achieves high performance in controlled settings but collapses under slight perturbation does not qualify as cognitively viable, regardless of benchmark success.

This framing explains why increasingly large models require extensive regularization, architectural constraints, and feedback mechanisms to remain stable.

14.2 Optimization Versus Viability

Much of modern machine learning is framed as optimization: minimizing loss, maximizing reward, or improving predictive accuracy. Cognitive Physics does not reject optimization, but it subordinates it to viability.

An optimizer that destabilizes itself through runaway gradients, feedback loops, or reward hacking fails the bounded disturbance condition. Conversely, systems that sacrifice optimality for robustness often perform better in open-ended environments.

This distinction clarifies why purely objective-driven agents frequently exhibit brittle or pathological behavior when deployed beyond their training regimes.

14.3 Generalization as Stability Across Distributions

Generalization is often treated as a statistical property: performance on unseen data drawn from a similar distribution. Cognitive Physics reframes generalization as stability across changing uncertainty regimes.

A system generalizes when its internal structure remains coherent despite shifts in input statistics. Distributional robustness, not accuracy alone, becomes the marker of cognitive viability.

This interpretation unifies generalization, robustness, and out-of-distribution performance under a single structural criterion.

14.4 Feedback, Memory, and Continual Learning

Artificial systems struggle with continual learning because new information often destabilizes existing structure. Catastrophic forgetting is not a technical oversight; it is a failure of stability preservation.

Cognitive Physics interprets memory not as storage, but as structural constraint. Persistent representations are those that can be re-entered without collapse under ongoing learning.

Successful continual learning architectures therefore resemble biological systems: they introduce modularity, replay, consolidation, and regulated plasticity to preserve coherence over time.

14.5 Scale and Architecture in AI

Modern AI systems operate across multiple scales: micro-level parameter updates, meso-level representations, and macro-level behaviors. Cognitive Physics insists that explanations of AI cognition specify the scale at which claims apply.

Architectures that ignore scale dependencies often misattribute failures to data scarcity or algorithm choice when the true issue is aggregation instability. Recognizing scale-indexed coherence allows more principled architectural design.

This perspective also explains why scaling laws improve performance only when accompanied by structural regularization.

14.6 Alignment as Viability, Not Obedience

Alignment is often framed as ensuring that artificial agents follow human intentions or values. Cognitive Physics reframes alignment more fundamentally: an aligned system is one that remains viable within the uncertainty of its operational environment without destabilizing itself or its surroundings.

Misalignment failures frequently arise from optimization without constraint, where systems pursue narrow objectives at the expense of broader coherence. Cognitive Physics treats alignment as a structural property, not a moral add-on.

An aligned system is one whose learning dynamics do not amplify harm, collapse coordination, or destroy the conditions of its own persistence.

14.7 Failure Modes in Artificial Cognition

Artificial systems exhibit well-documented collapse modes: gradient explosion, reward hacking, mode collapse, adversarial brittleness, and runaway feedback. Cognitive Physics interprets these as violations of bounded disturbance, connectivity viability, or scale coherence.

These failures are not incidental; they reveal the limits of admissible cognitive architectures. Any AI theory that cannot predict or explain these failures fails the collapse accountability condition.

14.8 What Artificial Intelligence Contributes to the Meeting Ground

Within the Cognitive Physics framework, artificial intelligence contributes:

- Explicit testbeds for cognitive viability
- Clear demonstrations of instability under unconstrained optimization
- Architectural strategies for preserving coherence
- Rapid experimental feedback on learning dynamics

AI does not redefine cognition; it exposes it. By failing loudly and repeatedly, artificial systems make visible the constraints that biological cognition has quietly satisfied all along.

Cognitive Physics hosts artificial intelligence as the fastest learner at the party—one that stumbles often, but teaches everyone else where the floor is slippery.

15 Ethics as Viability Under Constraint

Ethics is traditionally treated as a normative domain, separate from the descriptive sciences of cognition. Moral philosophy often asks what agents ought to do, while cognitive science asks how agents actually function. Cognitive Physics dissolves this separation by reframing ethics as a question of viability: which patterns of behavior, coordination, and decision-making allow cognitive systems to persist under uncertainty without collapsing themselves or their environments.

In this framework, ethics is not imposed from outside cognition. It emerges from the same structural constraints that govern learning, adaptation, and stability. Ethical failure is therefore not a violation of abstract rules, but a breakdown of persistence.

15.1 From Moral Rules to Stability Conditions

Many ethical systems are expressed as rules, duties, or values. While these formulations differ culturally and historically, Cognitive Physics asks what they have in common structurally. The answer is simple: ethical prescriptions tend to discourage behaviors that destabilize individuals, groups, or environments over time.

Violence, deception, exploitation, and unchecked short-term optimization often yield immediate gains but undermine long-term coherence. Cooperation, restraint, fairness, and reciprocity tend to stabilize interaction under uncertainty. Ethics, viewed through Cognitive Physics, encodes learned constraints discovered through repeated encounters with collapse.

This interpretation does not claim that ethical rules are universally correct. It claims that many ethical patterns persist because they support viability.

15.2 Ethical Behavior as Bounded Optimization

Ethical dilemmas are often framed as conflicts between maximizing outcomes and respecting constraints. Cognitive Physics resolves this tension by recognizing that unconstrained optimization is itself a source of instability.

Agents that pursue narrow objectives without regard for broader system coherence frequently generate cascading failures. Ethical behavior, in this sense, corresponds to bounded optimization: pursuing goals while maintaining conditions that allow continued interaction, trust, and learning.

This reframing aligns ethics with control theory and information theory. Just as a stable controller avoids aggressive gains that induce oscillation, a viable agent avoids actions that destabilize its social and ecological context.

15.3 Social Coherence and Shared Stability

Ethics becomes especially salient in multi-agent systems. Social environments introduce additional sources of uncertainty, including other agents with their own goals, learning dynamics, and failure modes.

Cognitive Physics treats ethical norms as mechanisms for regulating disturbance across agents. Norms constrain behavior in ways that reduce unpredictability, enable coordination, and preserve shared stability basins.

From this perspective, trust is not a moral virtue but a structural achievement. It emerges when agents reliably remain within bounds that others can predict and accommodate.

15.4 Scale Dependence in Ethical Reasoning

Ethical judgments are often sensitive to scale. Actions that appear harmless at one scale may be destructive at another. Cognitive Physics insists that ethical claims, like cognitive claims, must be indexed to the scale at which their effects manifest.

Short-term individual benefits may destabilize long-term collective viability. Local optimization may erode global coherence. Ethical reasoning that ignores scale risks promoting actions that succeed briefly but collapse over time.

Recognizing scale-indexed ethics prevents the elevation of narrow success metrics into universal principles.

15.5 Failure Modes: Ethical Collapse

Ethical collapse occurs when actions undermine the conditions that support continued interaction. This may manifest as erosion of trust, escalation of conflict, exploitation of shared resources, or breakdown of coordination.

Cognitive Physics treats such collapse not as moral failure in a metaphysical sense, but as a predictable outcome when bounded disturbance is violated in social systems. Ethical breakdown follows the same structural logic as cognitive or ecological collapse.

Any ethical framework that cannot specify how its principles fail under pressure fails the collapse accountability condition.

15.6 Ethics Without Moral Absolutes

Cognitive Physics does not require moral absolutes, universal values, or external authorities. It requires only that ethical patterns be evaluated by their contribution to persistence under uncertainty.

This does not imply relativism. Some behaviors reliably destabilize systems across contexts, while others reliably support coherence. Ethics becomes an empirical question: which norms sustain viable interaction across changing conditions?

In this sense, ethics is continuous with cognition and learning rather than opposed to them.

15.7 What Ethics Contributes to the Meeting Ground

Within the Cognitive Physics framework, ethics contributes:

- A constraint-based interpretation of moral norms
- A bridge between individual cognition and collective stability
- A non-mystical account of cooperation and trust
- Explicit failure modes for social collapse

Ethics is not an external judge at the cognitive party. It is the set of floor rules discovered through repeated falls. Systems that ignore them may dance briefly, but they do not remain standing.

16 Collective and Social Cognition

Cognition does not occur only within individual systems. Many of the most stable, adaptive, and powerful cognitive processes emerge at the collective level: groups, institutions, cultures, and societies routinely solve problems no individual could manage alone. Cognitive Physics treats collective cognition not as a metaphor or emergent curiosity, but as a direct extension of the same constraints that govern individual cognition.

The central claim is simple: collectives become cognitive when their interactions form stable, reconstructable patterns under uncertainty. When this condition is met, groups learn, remember, adapt, and fail in ways that closely mirror individual cognitive systems.

16.1 From Individual to Collective Stability

Individual cognition is bounded by biological limits: attention, memory, energy, and lifespan. Collective systems extend these limits by distributing cognitive labor across multiple agents and artifacts. Language, norms, tools, and institutions allow information to persist beyond any single individual.

Cognitive Physics interprets this distribution as an expansion of the stability basin. What matters is not where cognition resides, but whether the collective maintains coherence when agents enter, leave, disagree, or fail.

When coordination mechanisms succeed, the collective exhibits properties indistinguishable from cognition: shared beliefs, adaptive strategies, error correction, and learning over time.

16.2 Communication as Aggregation Mechanism

Communication is the primary mechanism by which individual cognitive states are aggregated into collective ones. Language, symbols, rituals, and signals act as observation kernels that filter, compress, and stabilize information across agents.

Cognitive Physics treats communication not as transmission of truth, but as regulation of uncertainty. Effective communication reduces ambiguity to manageable levels, enabling coordinated action. Ineffective communication amplifies disturbance, fragmenting collective coherence.

This framing explains why misinformation, noise, and overload destabilize groups even when individual agents remain cognitively intact.

16.3 Institutions as Memory and Control Structures

Institutions—legal systems, scientific norms, economic rules, educational structures—function as long-term memory and control mechanisms for collective cognition. They store constraints discovered through historical learning and enforce bounded behavior across generations.

From the perspective of Cognitive Physics, institutions are not arbitrary social constructs. They are structural responses to repeated collapse. Successful institutions persist because they stabilize interaction under uncertainty; failed institutions dissolve or are replaced.

This interpretation aligns institutional stability with the same principles governing neural memory and adaptive control.

16.4 Social Feedback and Error Correction

Collective cognition depends critically on feedback. Social feedback mechanisms—criticism, reputation, incentives, sanctions—regulate deviation and reinforce viable behavior.

Cognitive Physics treats social feedback as the collective analogue of neural error signals. When feedback is timely and proportional, collectives adapt smoothly. When feedback is delayed, distorted, or suppressed, instability grows.

This explains why echo chambers, authoritarian suppression, and unregulated amplification often precede social collapse: they disrupt bounded disturbance at the collective scale.

16.5 Scale and Timescale in Collective Cognition

Collective cognition operates on timescales far longer than individual cognition. Cultural norms may persist for centuries; scientific knowledge accumulates across generations. Cognitive Physics insists that these extended timescales be treated explicitly.

Short-term stability at the individual level can undermine long-term collective viability. Conversely, collective constraints may limit individual freedom to preserve coherence over time. Ethical and political tensions often arise from misalignment across scales rather than fundamental disagreement.

Recognizing scale-indexed coherence clarifies why local optimization can produce global failure.

16.6 Failure Modes: Social and Collective Collapse

Collective cognitive systems exhibit distinctive collapse modes: polarization, loss of trust, institutional decay, coordination breakdown, and violence. These failures follow the same structural logic as individual cognitive collapse.

When information flow becomes unbounded, feedback fails, or connectivity fragments, the collective loses reconstructability. Actions continue, but they no longer form a coherent trajectory.

Cognitive Physics treats social collapse as a diagnostic signal, not a moral mystery. It reveals where structural constraints have been violated.

16.7 What Collective Cognition Contributes to the Meeting Ground

Within the Cognitive Physics framework, collective and social cognition contribute:

- Demonstrations of cognition beyond individual agents
- Mechanisms for long-term memory and learning
- Scalable feedback and error correction systems
- Clear examples of multi-scale stability and collapse

They show that cognition is not confined to minds or machines. It is a property of interacting systems that manage uncertainty together.

Cognitive Physics hosts collective cognition as the largest ensemble on the dance floor—many bodies moving as one pattern, stable only as long as the shared rhythm holds.

17 Unification Without Finality

Scientific history is marked by repeated attempts to arrive at final theories—complete descriptions from which all phenomena could, in principle, be derived. While such efforts have driven enormous progress, they have repeatedly encountered a fundamental limitation: systems that learn, adapt, and persist in open environments do not admit final descriptions.

Cognitive Physics embraces this limitation explicitly. It proposes unification without finality: a framework that integrates diverse models of cognition without claiming to exhaust or terminate them. This stance is not a retreat from rigor, but a recognition of the structural nature of adaptive systems.

17.1 Why Final Theories Fail for Cognition

Final theories assume closed domains, fixed variables, and invariant laws. Cognition violates all three assumptions. Cognitive systems continually encounter novelty, modify their internal structure, and operate across shifting scales. Any theory that attempts to fully specify cognition at all times necessarily freezes it.

Cognitive Physics therefore rejects finality not on philosophical grounds, but on physical ones. A theory that cannot accommodate its own revision under new conditions fails the persistence criterion it seeks to explain.

This does not imply that cognition is unprincipled or arbitrary. It implies that principles govern how models change, not what they ultimately conclude.

17.2 Constraints Versus Content

The unifying power of Cognitive Physics lies in its separation of constraints from content. It does not dictate representations, mechanisms, or dynamics. Instead, it specifies the conditions under which any such choices remain viable.

Different cognitive models may disagree profoundly in content while agreeing on constraints. They may use different mathematics, substrates, or metaphors, yet still occupy the same admissible region of the cognitive landscape.

This separation allows unification without reduction and comparison without erasure.

17.3 Learning as an Open-Ended Process

Learning, whether biological, artificial, or collective, is inherently open-ended. New environments introduce new distinctions, new failure modes, and new stability requirements. Cognitive Physics treats learning as continuous basin reshaping rather than convergence to a final state.

This perspective aligns learning with evolution, control, and information theory. Systems do not learn to become finished; they learn to remain viable under changing conditions.

Any framework that treats learning as a path toward completion misunderstands its physical role.

17.4 Model Competition Without Ontological War

Scientific progress often degenerates into disputes over which model is fundamentally correct. Cognitive Physics reframes these disputes as competitions within a shared constraint space.

Models compete not for metaphysical dominance, but for viability: explanatory power, robustness, scalability, and predictive coherence under uncertainty. Multiple models may coexist if they occupy different scales or regimes.

This approach encourages pluralism without relativism and rigor without exclusion.

17.5 Revision as a Feature, Not a Bug

Because Cognitive Physics is itself a learning framework, it must remain revisable. New empirical results, new architectures, and new forms of cognition may require refinement or expansion of constraints.

This does not weaken the framework. It strengthens it. A framework that cannot update under evidence is structurally inconsistent with its own principles.

Unification without finality therefore demands humility as a methodological commitment.

17.6 What Unification Without Finality Contributes to the Meeting Ground

Within the Cognitive Physics framework, this stance contributes:

- Integration without reduction
- Coexistence without relativism
- Progress without closure
- Rigor without dogma

It ensures that the cognitive meeting ground remains open, adaptive, and scientifically honest—capable of hosting new ideas without collapsing into chaos or freezing into orthodoxy.

Cognitive Physics does not aim to end the conversation about cognition. It aims to keep the conversation standing.

18 Methodological Implications, Predictions, and Falsifiability

A framework that claims to unify cognition across scales must make clear how it can fail. Without explicit methodological commitments and testable consequences, Cognitive

Physics would risk becoming a descriptive umbrella rather than a scientific discipline. This section therefore articulates how Cognitive Physics constrains inquiry, generates predictions, and exposes itself to falsification.

Cognitive Physics does not compete with existing cognitive sciences by replacing their models. It competes by imposing structural discipline: models that violate persistence under uncertainty are rejected regardless of their explanatory appeal.

18.1 Methodological Commitments

Cognitive Physics rests on four non-negotiable methodological commitments.

First, cognition must be treated as a physical phenomenon. Even when implemented symbolically or socially, cognitive processes consume resources, operate under noise, and unfold in time. Any model that abstracts away physical constraints without justification is incomplete.

Second, explanations must be scale-explicit. Claims about cognition must specify the spatial, temporal, and organizational scale at which they apply. Apparent contradictions between models often dissolve when scale dependence is made explicit.

Third, stability takes precedence over optimality. Cognitive Physics evaluates models by whether they remain coherent under perturbation, not by whether they maximize a chosen objective. Optimization is admissible only insofar as it preserves viability.

Fourth, collapse is informative. Breakdowns of cognition—whether neural, artificial, or social—are not anomalies to be ignored. They are essential data that reveal the boundaries of admissible structure.

These commitments collectively define what it means to do Cognitive Physics as a scientific practice.

18.2 Predictions at the Structural Level

Cognitive Physics does not predict specific thoughts, behaviors, or beliefs. Instead, it generates structural predictions that apply across domains.

One such prediction is the existence of sharp transitions between cognitive coherence and collapse as disturbance exceeds regulatory capacity. These transitions should appear as non-linear breakdowns rather than gradual degradation, regardless of substrate.

Another prediction is that learning systems which improve performance by relaxing stability constraints will eventually exhibit catastrophic failure. Short-term gains achieved through unconstrained optimization should correlate with long-term brittleness.

A third prediction is that viable cognitive systems will exhibit identifiable bottlenecks in information flow. Systems that attempt to process all available information without compression should fail under realistic conditions.

These predictions are qualitative but falsifiable. Their absence across domains would undermine the framework.

18.3 Predictions Across Domains

In neuroscience, Cognitive Physics predicts that loss of recurrent connectivity or feedback regulation will precede cognitive collapse even if local neural activity remains intact. Disorders of consciousness, attention, and coordination should correlate more strongly with connectivity disruption than with localized damage alone.

In artificial intelligence, the framework predicts that systems optimized aggressively for narrow objectives will display instability under distributional shift, adversarial input, or long-horizon deployment. Robust systems should exhibit explicit mechanisms for bounding disturbance, even at the cost of reduced peak performance.

In collective systems, Cognitive Physics predicts that breakdowns in communication and feedback will precede social collapse. Polarization, institutional decay, and coordination failure should follow measurable increases in unbounded information flow and loss of shared constraints.

These predictions do not depend on specific mechanisms. They depend only on whether persistence under uncertainty is preserved.

18.4 Falsification Criteria

Cognitive Physics is falsified if any of the following are demonstrated:

- A cognitive system that persists indefinitely under uncertainty without regulating disturbance.
- A learning system that remains stable while allowing unbounded amplification of error or noise.
- A cognitive architecture that generalizes robustly without bottlenecks, feedback, or constraint.

- A collective system that maintains coherence without communication, norms, or regulatory structure.

The discovery of such a system would directly contradict the core claim that persistence requires constraint.

Importantly, improved performance alone does not falsify the framework. Only violation of structural necessity does.

18.5 Experimental and Computational Pathways

Cognitive Physics invites testing through multiple methodologies. In neuroscience, this includes perturbation experiments that selectively disrupt connectivity while preserving local activity. In AI, this includes stress-testing learning systems under distributional shift, delayed feedback, and constrained resources. In social systems, this includes longitudinal analysis of communication patterns and institutional stability.

Because the framework is scale-agnostic, confirmation or refutation may occur in any domain. No single experiment can validate Cognitive Physics, but a single decisive counterexample can invalidate it.

18.6 Risk and Scientific Discipline

By design, Cognitive Physics accepts scientific risk. It does not shield itself behind interpretive flexibility or metaphysical claims. If persistence under uncertainty proves insufficient to constrain cognition, the framework must be revised or abandoned.

This vulnerability is not a weakness. It is the defining feature of a scientific meeting ground rather than a belief system.

Cognitive Physics stakes its claim on a simple but demanding proposition: that cognition, wherever it appears, cannot escape the constraints of stability, scale, and failure.

19 Conclusion: Cognition as Persistent Structure

This paper has advanced a simple but demanding claim: cognition is not defined by intelligence, consciousness, representation, or optimization, but by persistence under uncertainty. Wherever cognition appears—neuronal, artificial, embodied, collective—it does so as a structure that maintains coherence while navigating disturbance, noise, and change.

Cognitive Physics does not propose a new object to study. It proposes a new criterion for admissibility. Models of cognition are not evaluated by elegance, popularity, or metaphysical depth, but by whether they can remain standing when uncertainty is sustained.

Across neuroscience, Bayesian inference, control theory, information theory, evolution, embodiment, artificial intelligence, ethics, and collective systems, the same pattern recurs. Cognition survives by bounding disturbance, regulating information flow, sustaining viable connectivity, and acknowledging its own failure modes. These are not domain-specific tricks; they are structural necessities.

Importantly, Cognitive Physics does not seek to replace existing sciences of cognition. It hosts them. It provides a shared meeting ground where diverse models can coexist, compete, and evolve—so long as they respect the non-negotiable constraints imposed by reality itself. The framework unifies without finality, integrates without reduction, and advances without closure.

Because cognition is adaptive, the framework that studies it must be adaptive as well. Cognitive Physics remains open to revision, expansion, and falsification. Its success will not be measured by agreement, but by its ability to clarify why certain cognitive systems persist while others collapse.

If this framework fails, it should fail clearly. If it holds, it will not end inquiry, but sharpen it—turning questions of mind, intelligence, and behavior into questions of structure, stability, and survival.

Cognition, in this view, is not a thing that acts upon the world. It is a pattern that continues.

20 Conclusion

Cognitive Physics does not seek to end debate about cognition. It seeks to give debate a shared ground where models can meet, interact, and evolve without collapsing into relativism or dogma. By grounding cognition in persistence under uncertainty, the framework allows science to remain open-ended, self-correcting, and structurally disciplined.

This is not a final theory of mind. It is a theory of how theories of mind survive.

References

[1] Marr, D. (1982). *Vision: A Computational Investigation into the Human Representation and Processing of Visual Information*. San Francisco: W. H. Freeman.

[2] Newell, A., & Simon, H. A. (1976). Computer science as empirical inquiry: Symbols and search. *Communications of the ACM*, 19(3), 113–126.

[3] Kuhn, T. S. (1962). *The Structure of Scientific Revolutions*. Chicago: University of Chicago Press.

[4] Sporns, O. (2011). *Networks of the Brain*. Cambridge, MA: MIT Press.

[5] Friston, K. (2010). The free-energy principle: A unified brain theory? *Nature Reviews Neuroscience*, 11(2), 127–138.

[6] Breakspear, M. (2017). Dynamic models of large-scale brain activity. *Nature Neuroscience*, 20(3), 340–352.

[7] Ogata, K. (2010). *Modern Control Engineering* (5th ed.). Upper Saddle River, NJ: Prentice Hall.

[8] Khalil, H. K. (2002). *Nonlinear Systems* (3rd ed.). Upper Saddle River, NJ: Prentice Hall.

[9] Åström, K. J., & Murray, R. M. (2010). *Feedback Systems: An Introduction for Scientists and Engineers*. Princeton, NJ: Princeton University Press.

[10] Shannon, C. E. (1948). A mathematical theory of communication. *Bell System Technical Journal*, 27, 379–423, 623–656.

[11] Cover, T. M., & Thomas, J. A. (2006). *Elements of Information Theory* (2nd ed.). New York: Wiley.

[12] Tishby, N., Pereira, F. C., & Bialek, W. (2000). The information bottleneck method. *arXiv:physics/0004057*.

[13] Maynard Smith, J. (1982). *Evolution and the Theory of Games*. Cambridge: Cambridge University Press.

[14] Gould, S. J. (2002). *The Structure of Evolutionary Theory*. Cambridge, MA: Harvard University Press.

[15] Kirschner, M., & Gerhart, J. (1998). Evolvability. *Proceedings of the National Academy of Sciences*, 95(15), 8420–8427.

[16] Varela, F. J., Thompson, E., & Rosch, E. (1991). *The Embodied Mind*. Cambridge, MA: MIT Press.

[17] Clark, A. (1997). *Being There: Putting Brain, Body, and World Together Again*. Cambridge, MA: MIT Press.

[18] Chemero, A. (2009). *Radical Embodied Cognitive Science*. Cambridge, MA: MIT Press.

[19] Sutton, R. S., & Barto, A. G. (2018). *Reinforcement Learning: An Introduction* (2nd ed.). Cambridge, MA: MIT Press.

[20] Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. Cambridge, MA: MIT Press.

[21] French, R. M. (1999). Catastrophic forgetting in connectionist networks. *Trends in Cognitive Sciences*, 3(4), 128–135.

[22] Hutchins, E. (1995). *Cognition in the Wild*. Cambridge, MA: MIT Press.

[23] Helbing, D. (2013). Globally networked risks and how to respond. *Nature*, 497(7447), 51–59.

[24] Ostrom, E. (2009). A general framework for analyzing sustainability of social-ecological systems. *Science*, 325(5939), 419–422.

[25] Ashby, W. R. (1956). *An Introduction to Cybernetics*. London: Chapman & Hall.

[26] Jonas, H. (1984). *The Imperative of Responsibility*. Chicago: University of Chicago Press.

[27] Holland, J. H. (1992). Complex adaptive systems. *Daedalus*, 121(1), 17–30.

[28] Mitchell, M. (2009). *Complexity: A Guided Tour*. Oxford: Oxford University Press.